

STAT 4710 Final Project -

Andrew Raine, Graham Branscom, Mahika Calyanakoti

Part 1: Imports and Loading the Dataset

[] ↪ 7 cells hidden

Part 2: EDA

Data cleaning and feature engineering

```
1 #@title Data cleaning and feature engineering
2 # Checking for nans; none exist
3 na_counts = df.isna().sum()
4
5 # filtering out crab outliers
6 df = df[df.Height < 1]
7
8 # binarizing the age
9 median_age = df['Age'].median()
10 df['Young|Old'] = (df['Age'] > median_age).astype(int)

1 df.to_csv('cleaned_data.csv', index=False)
```

Crab summary

Show code

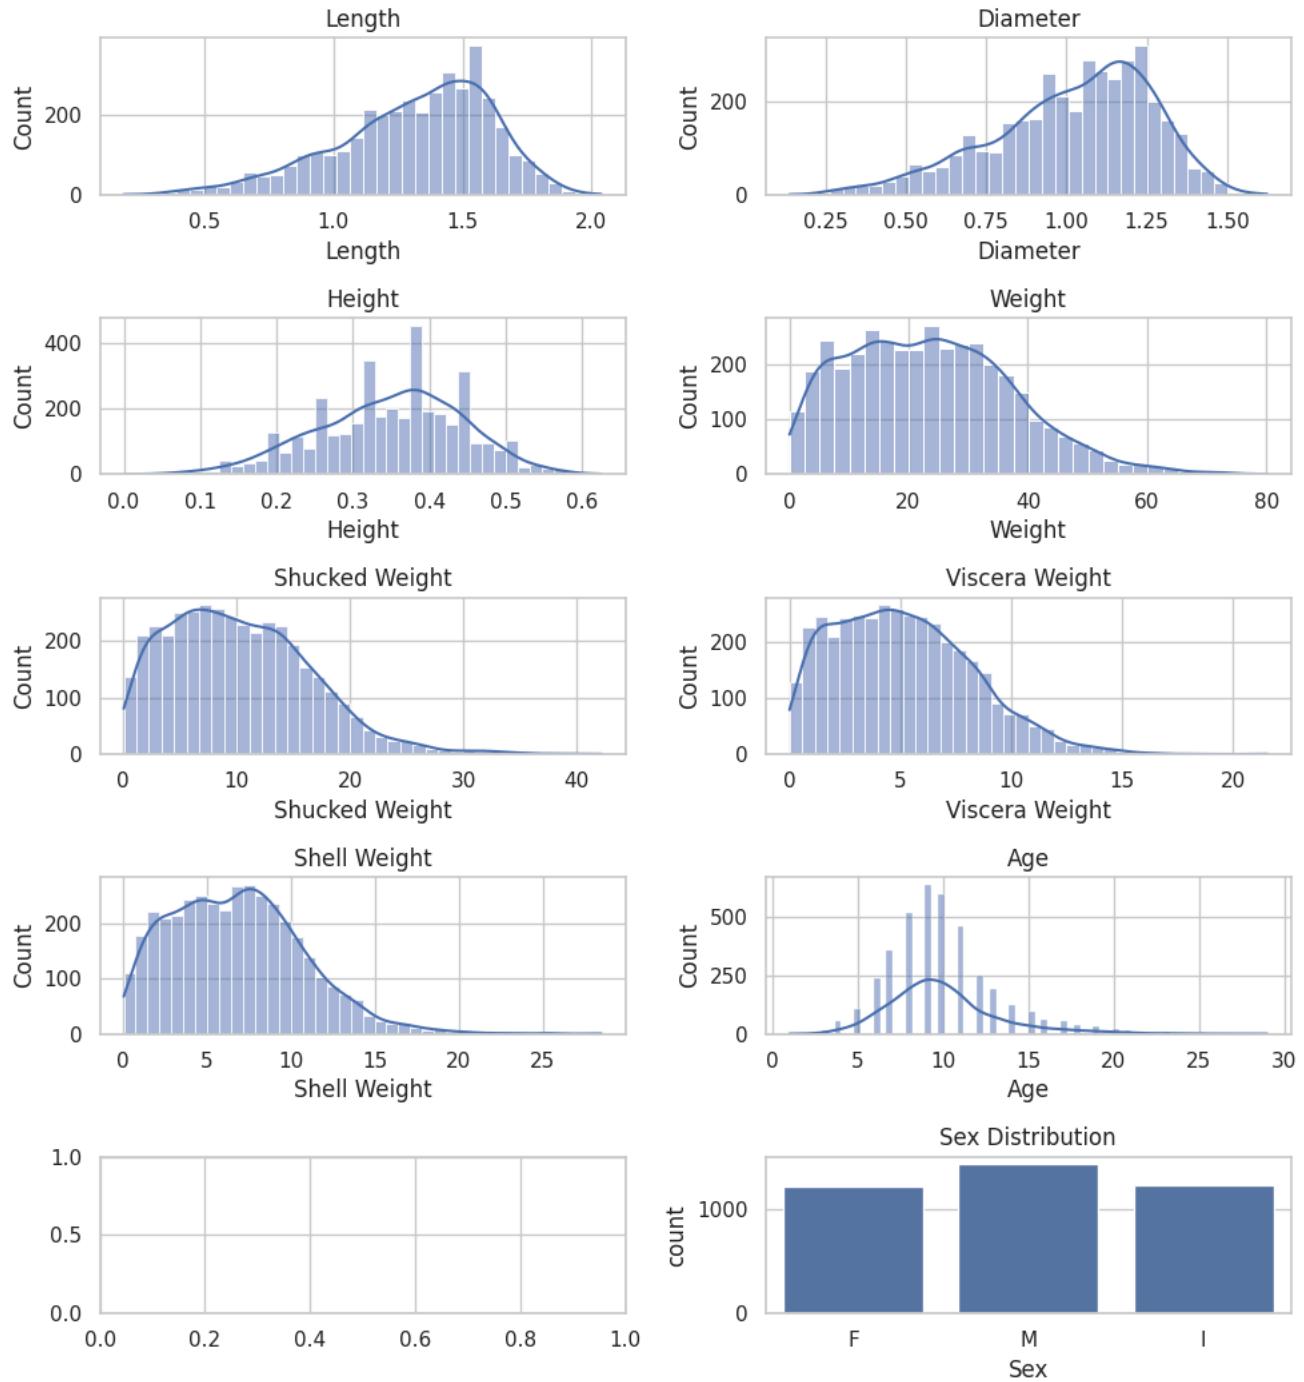
	Length	Diameter	Height	Weight	Shucked Weight	Viscera Weight	Shell Weight	Age	Young Old
count	3891.000000	3891.000000	3891.000000	3891.000000	3891.000000	3891.000000	3891.000000	3891.000000	3891.000000
mean	1.311234	1.020827	0.348497	23.558959	10.202101	5.134797	6.794634	9.955281	0.350039
std	0.300408	0.248208	0.096037	13.880211	6.267671	3.101686	3.942167	3.221642	0.477043
min	0.187500	0.137500	0.000000	0.056699	0.028349	0.014175	0.042524	1.000000	0.000000
25%	1.125000	0.875000	0.287500	12.672227	5.336793	2.664853	3.713785	8.000000	0.000000
50%	1.362500	1.062500	0.362500	22.792998	9.539607	4.861939	6.662133	10.000000	0.000000
75%	1.537500	1.200000	0.412500	32.786197	14.266886	7.200773	9.355335	11.000000	1.000000
max	2.037500	1.625000	0.625000	80.101512	42.184056	21.545620	28.491248	29.000000	1.000000

Variable distributions

Show code




Distribution of Variables



✓ PCA

```
1 #@title PCA
2 from sklearn.decomposition import PCA
3
4 # Independent variables for PCA (excluding 'Age' and 'Young|Old')
5 pca_vars = ['Length', 'Diameter', 'Height', 'Weight', 'Shucked Weight', 'Viscera Weight', 'Shell Weight']
6 X_pca = df[pca_vars]
7
8 # Standardizing the data for PCA
9 scaler_pca = StandardScaler()
10 X_pca_scaled = scaler_pca.fit_transform(X_pca)
11
12 # Applying PCA to reduce dimensions to 2
13 pca = PCA(n_components=2)
14 principal_components = pca.fit_transform(X_pca_scaled)
15
16 # Creating a DataFrame for the principal components
17 pca_df = pd.DataFrame(data=principal_components, columns=['PC1', 'PC2'])
18 # Including the 'Young|Old' variable for color coding
19 pca_df['Young|Old'] = df['Young|Old']
20
21 # Plotting the 2D scatter plot of the two principal components with color coding
22 plt.figure(figsize=(8, 6))
23 sns.scatterplot(x='PC1', y='PC2', hue='Young|Old', data=pca_df, alpha=0.6)
24 plt.title('PCA of Crab Variables')
25 plt.xlabel('Principal Component 1')
26 plt.ylabel('Principal Component 2')
27 plt.grid(True)
28 plt.show()
29
30
31 print("\n\nPCA Loadings")
32 for pc_idx in range(2):
33     for var, loading in list(zip(pca_vars, pca.components_[pc_idx])):
34         print('\t' + var + ":", loading)
35     print('\n')
```



PCA Loadings

Length: 0.3806276567338248
 Diameter: 0.3810108853408206
 Height: 0.36710538729474895
 Weight: 0.3874789211979131
 Shucked Weight: 0.37480648508033465
 Viscera Weight: 0.3782665434524006
 Shell Weight: 0.3761403106207747

Length: -0.3043014272247193
 Diameter: -0.3390792470684685
 Height: -0.5624537187565186
 Weight: 0.3132958052943968
 Shucked Weight: 0.4984521051456185
 Viscera Weight: 0.3592745479043055
 Shell Weight: 0.019615206918008693

3D PCA of Crab Variables Colored By Young|Old


```

1 #@title 3D PCA of Crab Variables Colored By Young|Old
2 # Perform PCA to reduce to 3 dimensions for visualization
3 pca = PCA(n_components=3)
4 reduced_embeddings = pca.fit_transform(X_pca_scaled)
5
6 # Create an interactive 3D scatter plot
7 fig = go.Figure(data=[go.Scatter3d(
8     x=reduced_embeddings[:, 0],
9     y=reduced_embeddings[:, 1],
10    z=reduced_embeddings[:, 2],
11    mode='markers',
12    marker=dict(
13        size=5,
14        color=df['Young|Old'],          # set color to an array/list of desired values
15        colorscale='Viridis',          # choose a colorscale
16        opacity=0.8
17    )
18 )])
19
20 # Update the layout to add titles and axis labels
21 fig.update_layout(
22     title='3D PCA of Crab Variables',
23     scene=dict(
24         xaxis_title='PC 1',
25         yaxis_title='PC 2',
26         zaxis_title='PC 3'
27     ),
28     margin=dict(r=10, l=10, b=10, t=30)
29 )
30
31 # Show the plot
32 fig.show()

```



3D PCA of Crab Variables




> 3D PCA of Crab Variables Colored By Age


[Show code](#)



3D PCA of Crab Variables



> t-SNE plot

[Show code](#)

✓ 3D t-SNE plot


```

1 #@title 3D t-SNE plot
2
3 from sklearn.manifold import TSNE
4 import plotly.graph_objs as go
5 import pandas as pd
6
7 # Assuming 'df' is your DataFrame and 'X_pca_scaled' is your scaled feature set.
8
9 # Setting up t-SNE with 3 components for a 3D visualization
10 tsne_3d = TSNE(n_components=3, random_state=42)
11
12 # Applying t-SNE to the scaled data
13 tsne_results_3d = tsne_3d.fit_transform(X_pca_scaled)
14
15 # Creating a DataFrame for the t-SNE components
16 tsne_df_3d = pd.DataFrame(data=tsne_results_3d, columns=['TSNE1', 'TSNE2', 'TSNE3'])
17 tsne_df_3d['Young|Old'] = df['Young|Old']
18
19 # Create an interactive 3D scatter plot using plotly
20 fig = go.Figure(data=[go.Scatter3d(
21     x=tsne_df_3d['TSNE1'],
22     y=tsne_df_3d['TSNE2'],
23     z=tsne_df_3d['TSNE3'],
24     mode='markers',
25     marker=dict(
26         size=5,
27         color=tsne_df_3d['Young|Old'], # Color by 'Young|Old' binary class
28         colorscale='Viridis',         # Choose a colorscale
29         opacity=0.8
30     )
31 )])
32
33 # Update the layout to add titles and axis labels
34 fig.update_layout(
35     title='3D t-SNE Visualization of Crab Variables',
36     scene=dict(
37         xaxis_title='t-SNE Component 1',
38         yaxis_title='t-SNE Component 2',
39         zaxis_title='t-SNE Component 3'
40     ),
41     margin=dict(r=10, l=10, b=10, t=30)
42 )
43
44 # Show the plot
45 fig.show()
46

```



3D t-SNE Visualization of Crab Variables



Part 3: Linear Models

[] ↪ 5 cells hidden

Part 4: Tree Models

Random Forest Regressor

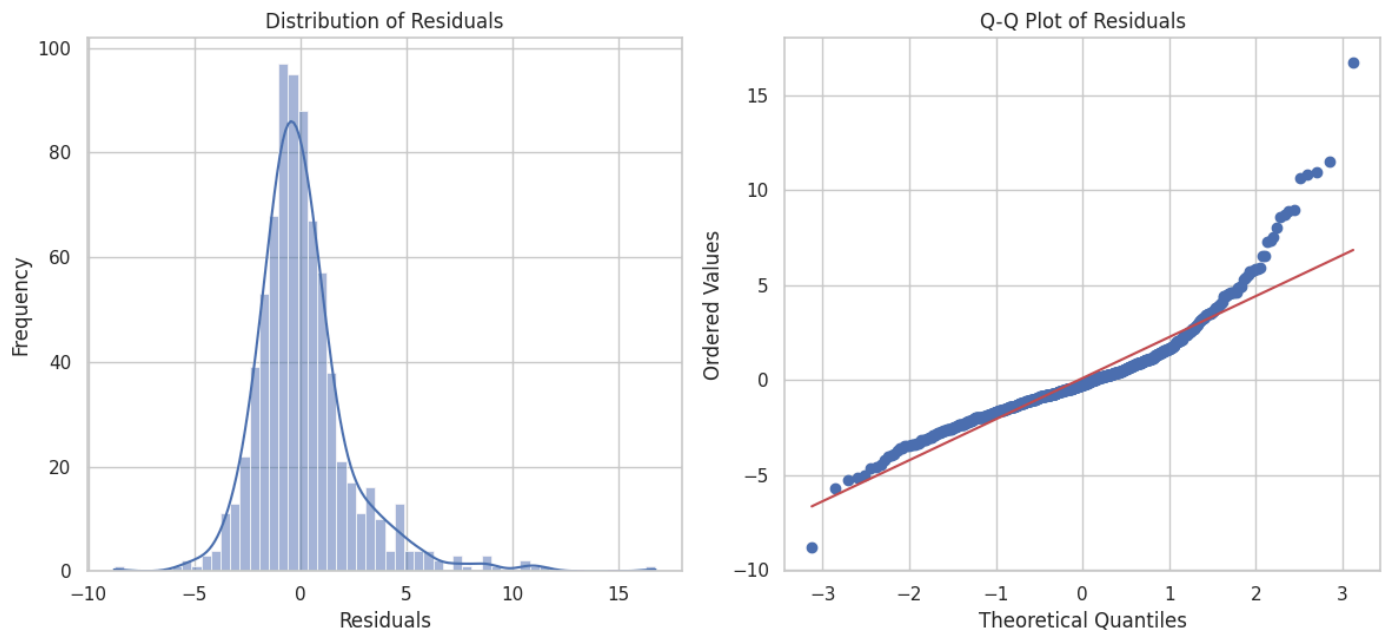
```

1 #@title Random Forest Regressor
2 from sklearn.ensemble import RandomForestRegressor
3 from sklearn.metrics import mean_squared_error
4 from scipy.stats import probplot
5 from sklearn.metrics import r2_score
6
7 # Selecting predictor variables (excluding 'Young|Old')
8 predictor_vars = ['Length', 'Diameter', 'Height', 'Weight', 'Shucked Weight', 'Viscera Weight', 'Shell Weight']
9 X_rf = df[predictor_vars]
10 y_rf = df['Age']
11
12 # Splitting the data into training and testing sets
13 X_train_rf, X_test_rf, y_train_rf, y_test_rf = train_test_split(X_rf, y_rf, test_size=0.2, random_state=42)
14
15 # Creating a random forest regressor model
16 random_forest_model = RandomForestRegressor(n_estimators=100, random_state=42)
17
18 # Fitting the model
19 random_forest_model.fit(X_train_rf, y_train_rf)
20
21 # Predicting on the test data
22 y_pred_rf = random_forest_model.predict(X_test_rf)
23
24 # Calculating residuals
25 residuals_rf = y_test_rf - y_pred_rf
26
27 # R^2 score
28 r2_rf = r2_score(y_test_rf, y_pred_rf)
29 print("R2:", r2_rf)
30
31 # Mean Squared Error
32 mse_rf = mean_squared_error(y_test_rf, y_pred_rf)
33 print("MSE:", mse_rf, "\n\n")
34
35
36 # Creating a figure to hold subplots
37 fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(12, 6))
38 fig.suptitle('Residual Analysis', fontsize=16)
39
40 # Plotting the histogram of residuals
41 sns.histplot(residuals_rf, kde=True, ax=axes[0])
42 axes[0].set_title('Distribution of Residuals')
43 axes[0].set_xlabel('Residuals')
44 axes[0].set_ylabel('Frequency')
45
46 # Creating Q-Q plot
47 probplot(residuals_rf, dist="norm", plot=axes[1])
48 axes[1].set_title('Q-Q Plot of Residuals')
49 axes[1].set_xlabel('Theoretical Quantiles')
50 axes[1].set_ylabel('Ordered Values')
51
52 plt.tight_layout()
53 plt.show()

```


R2: 0.5220396296821315
MSE: 5.233945699614891


Residual Analysis



```

1 # Assuming best_rf_model is your trained RandomForestClassifier or RandomForestRegressor
2 feature_importances = random_forest_model.feature_importances_
3 feature_importances_sorted = sorted(zip(predictor_vars, feature_importances), key=lambda x: x[1], reverse=False)
4 sorted_feature_names, sorted_feature_importances = zip(*feature_importances_sorted)
5
6 plt.figure(figsize=(10, 6))
7 plt.barh(sorted_feature_names, sorted_feature_importances, color='skyblue')
8 # plt.barh(predictor_vars, feature_importances_sorted, color='skyblue')
9 plt.xlabel('Feature Importance')
10 plt.ylabel('Feature')
11 plt.title('Feature Importance from Random Forest Model')
12 plt.show()

```



> XGBoost Regressor

[Show code](#)



R2: 0.4689048606432701
MSE: 5.815802508634486


The random forest regressor is the best model yet, even beating the XGBoost regressor.

✓ Random Forest Classification

```

1 #@title Random Forest Classification
2 from sklearn.ensemble import RandomForestClassifier
3 from sklearn.metrics import accuracy_score, precision_score, recall_score, roc_auc_score
4
5 # Assuming 'Young|Old' is binary and properly encoded, i.e., 0 for 'Young' and 1 for 'Old'
6 # Prepare the data
7 y_classifier = df['Young|Old']
8
9 # Split the data into training and testing sets
10 X_train_class, X_test_class, y_train_class, y_test_class = train_test_split(X_rf, y_classifier, test_size=0.2, random_state=42)
11
12 # Create a Random Forest Classifier model
13 random_forest_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
14
15 # Fit the model
16 random_forest_classifier.fit(X_train_class, y_train_class)
17
18 # Predict on the test data
19 y_pred_class = random_forest_classifier.predict(X_test_class)
20 y_pred_proba_class = random_forest_classifier.predict_proba(X_test_class)[:, 1] # probabilities for the positive class
21
22 # Evaluate the model
23 accuracy_class = accuracy_score(y_test_class, y_pred_class)
24 precision_class = precision_score(y_test_class, y_pred_class)
25 recall_class = recall_score(y_test_class, y_pred_class)
26 roc_auc_class = roc_auc_score(y_test_class, y_pred_proba_class)
27
28 # Residual-like analysis: Difference between actual and predicted probabilities
29 residuals_class = y_test_class - y_pred_proba_class
30
31 print("Accuracy:", accuracy_class)
32 print("Precision:", precision_class)
33 print("Recall:", recall_class)
34 print("ROC AUC:", roc_auc_class)

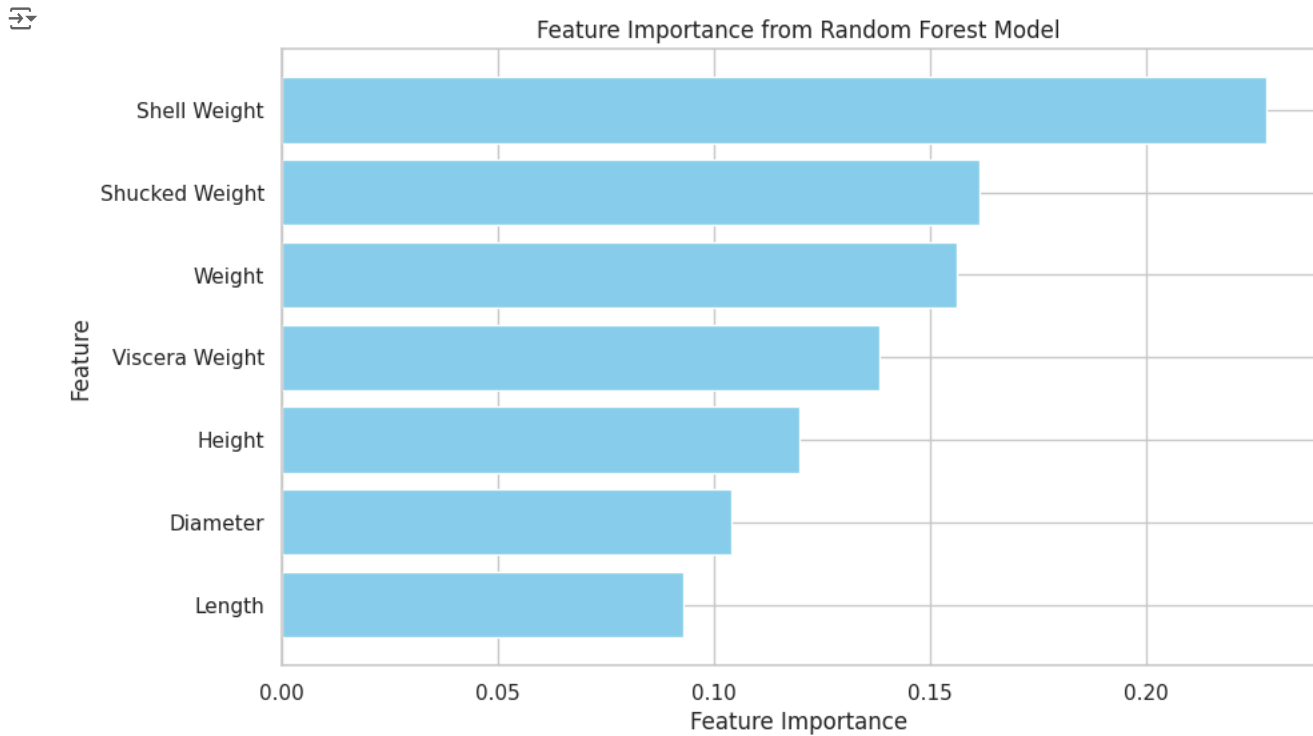
```

 Accuracy: 0.7753530166880617
 Precision: 0.73568281938326
 Recall: 0.5921985815602837
 ROC AUC: 0.8449277223625441

```

1 # Assuming best_rf_model is your trained RandomForestClassifier or RandomForestRegressor
2 feature_importances = random_forest_classifier.feature_importances_
3 feature_importances_sorted = sorted(zip(predictor_vars, feature_importances), key=lambda x: x[1], reverse=False)
4 sorted_feature_names, sorted_feature_importances = zip(*feature_importances_sorted)
5
6 plt.figure(figsize=(10, 6))
7 plt.barh(sorted_feature_names, sorted_feature_importances, color='skyblue')
8 # plt.barh(predictor_vars, feature_importances_sorted, color='skyblue')
9 plt.xlabel('Feature Importance')
10 plt.ylabel('Feature')
11 plt.title('Feature Importance from Random Forest Model')
12 plt.show()

```



✕ XGBoost Classification

```

1 #@title XGBoost Classification
2 # Prepare the data for classification
3 y_class = df['Young|Old']
4 X_train_class, X_test_class, y_train_class, y_test_class = train_test_split(X_rf, y_class, test_size=0.2, random_state=42)
5
6 # Create the XGBoost classifier model
7 xgb_classifier = xgb.XGBClassifier(objective='binary:logistic', n_estimators=100, seed=42, use_label_encoder=False, eval_metric='logloss')
8
9 # Fit the model
10 xgb_classifier.fit(X_train_class, y_train_class)
11
12 # Predict on the test data
13 y_pred_class = xgb_classifier.predict(X_test_class)
14 y_pred_proba_class = xgb_classifier.predict_proba(X_test_class)[:, 1] # probabilities for the positive class
15
16 # Evaluate the model
17 accuracy_class = accuracy_score(y_test_class, y_pred_class)
18 precision_class = precision_score(y_test_class, y_pred_class)
19 recall_class = recall_score(y_test_class, y_pred_class)
20 roc_auc_class = roc_auc_score(y_test_class, y_pred_proba_class)
21
22 print("Accuracy:", accuracy_class)
23 print("Precision:", precision_class)
24 print("Recall:", recall_class)
25 print("ROC AUC:", roc_auc_class)

```

```

Accuracy: 0.7573812580231065
Precision: 0.6882591093117408
Recall: 0.6028368794326241
ROC AUC: 0.834375044593804

```

✕ Part 5: Artificial Neural Network

> Data preparation

[Show code](#)

<ipython-input-126-57f296e4ce7b>:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view

› DataLoaders

[Show code](#)

▼ Model

```
1 #@title Model
2 import torch.nn as nn
3
4 class FFN(nn.Module):
5     def __init__(self, task):
6         super(FFN, self).__init__()
7
8         assert task in {"regression", "classification"}
9
10        self.task = task
11
12        self.backbone = nn.ModuleList([nn.BatchNorm1d(8, affine=False)])
13        for i in range(5):
14            self.backbone.append(nn.Linear(8, 8))
15            self.backbone.append(nn.Tanh())
16            # self.backbone.append(nn.Dropout(0.2))
17        self.backbone.append(nn.Linear(8, 4))
18
19        if task == "regression":
20            self.head = nn.Linear(4, 1)
21        else:
22            self.head = nn.Linear(4, 2)
23
24
25        def forward(self, x, raw=False):
26            for layer in self.backbone:
27                x = layer(x)
28            x_logits = x
29            x = self.head(x_logits)
30
31            if self.task == "classification":
32                x = nn.Softmax(dim=1)(x)
33
34            if raw:
35                return x, x_logits
36            return x
```

› Loss function and optimizer

[Show code](#)

› Train and test functions

[Show code](#)

▼ Classification

```
1 #@title Classification
2 class_model, class_test_data = trainAndTest("classification")
```

```
↩ Epoch 1
Batch 10: Loss = 0.6867, Accuracy = 65.00%
Batch 20: Loss = 0.6746, Accuracy = 69.69%
Batch 30: Loss = 0.6700, Accuracy = 66.25%
Batch 40: Loss = 0.6583, Accuracy = 66.56%
Batch 50: Loss = 0.6425, Accuracy = 69.06%
Batch 60: Loss = 0.6570, Accuracy = 58.12%
Batch 70: Loss = 0.6247, Accuracy = 66.88%
Batch 80: Loss = 0.6158, Accuracy = 72.50%
```

Batch 90: Loss = 0.6129, Accuracy = 69.69%
 Test Loss: 0.5949, Accuracy: 73.04%

Epoch 2

Batch 10: Loss = 0.5915, Accuracy = 69.69%
 Batch 20: Loss = 0.6128, Accuracy = 67.19%
 Batch 30: Loss = 0.5931, Accuracy = 70.31%
 Batch 40: Loss = 0.5794, Accuracy = 72.81%
 Batch 50: Loss = 0.5591, Accuracy = 74.69%
 Batch 60: Loss = 0.5724, Accuracy = 70.94%
 Batch 70: Loss = 0.5543, Accuracy = 73.44%
 Batch 80: Loss = 0.5917, Accuracy = 70.31%
 Batch 90: Loss = 0.5777, Accuracy = 70.31%
 Test Loss: 0.5659, Accuracy: 74.58%

Epoch 3

Batch 10: Loss = 0.5434, Accuracy = 75.62%
 Batch 20: Loss = 0.5478, Accuracy = 75.31%
 Batch 30: Loss = 0.5965, Accuracy = 69.69%
 Batch 40: Loss = 0.5666, Accuracy = 73.44%
 Batch 50: Loss = 0.5483, Accuracy = 74.69%
 Batch 60: Loss = 0.5623, Accuracy = 73.12%
 Batch 70: Loss = 0.5544, Accuracy = 74.06%
 Batch 80: Loss = 0.5720, Accuracy = 72.19%
 Batch 90: Loss = 0.5334, Accuracy = 77.19%
 Test Loss: 0.5465, Accuracy: 76.51%

Epoch 4

Batch 10: Loss = 0.5283, Accuracy = 78.12%
 Batch 20: Loss = 0.5423, Accuracy = 75.94%
 Batch 30: Loss = 0.5279, Accuracy = 77.50%
 Batch 40: Loss = 0.5765, Accuracy = 71.25%
 Batch 50: Loss = 0.5404, Accuracy = 75.94%
 Batch 60: Loss = 0.5561, Accuracy = 73.44%
 Batch 70: Loss = 0.5458, Accuracy = 75.31%
 Batch 80: Loss = 0.5306, Accuracy = 76.88%
 Batch 90: Loss = 0.5455, Accuracy = 76.25%
 Test Loss: 0.5369, Accuracy: 76.89%

Epoch 5

Batch 10: Loss = 0.5460, Accuracy = 75.62%
 Batch 20: Loss = 0.5245, Accuracy = 76.88%
 Batch 30: Loss = 0.5317, Accuracy = 76.56%
 Batch 40: Loss = 0.5309, Accuracy = 77.81%
 Batch 50: Loss = 0.5264, Accuracy = 76.88%

▼ Regression

```
1 #@title Regression
2 reg_model, reg_test_data = trainAndTest("regression", 50)
```



Epoch 1

Batch 10: Loss = 91.1427
 Batch 20: Loss = 91.2421
 Batch 30: Loss = 94.0418
 Batch 40: Loss = 90.6227
 Batch 50: Loss = 87.9925
 Batch 60: Loss = 79.3302
 Batch 70: Loss = 82.3218
 Batch 80: Loss = 71.0317
 Batch 90: Loss = 72.8962
 Test Loss: 58.3384

Epoch 2

Batch 10: Loss = 58.8409
 Batch 20: Loss = 60.1046
 Batch 30: Loss = 42.6817
 Batch 40: Loss = 40.2391
 Batch 50: Loss = 35.5226
 Batch 60: Loss = 29.0844
 Batch 70: Loss = 28.8568
 Batch 80: Loss = 29.2687
 Batch 90: Loss = 24.2316
 Test Loss: 21.1164

Epoch 3

Batch 10: Loss = 20.4493

```

Batch 20: Loss = 15.2075
Batch 30: Loss = 14.2824
Batch 40: Loss = 15.0294
Batch 50: Loss = 14.6226
Batch 60: Loss = 12.1941
Batch 70: Loss = 15.1486
Batch 80: Loss = 12.5043
Batch 90: Loss = 11.0391
Test Loss: 10.0775

```

```

Epoch 4
Batch 10: Loss = 10.3545
Batch 20: Loss = 9.0284
Batch 30: Loss = 9.5050
Batch 40: Loss = 9.2756
Batch 50: Loss = 10.5331
Batch 60: Loss = 7.9792
Batch 70: Loss = 7.2549
Batch 80: Loss = 9.2489
Batch 90: Loss = 8.0763
Test Loss: 7.8557

```

```

Epoch 5
Batch 10: Loss = 7.9468
Batch 20: Loss = 5.3029
Batch 30: Loss = 6.9176
Batch 40: Loss = 7.0757
Batch 50: Loss = 8.2167


```

✓ PCA on the neural network regression latent vectors

```

1 #@title PCA on the neural network regression latent vectors
2 embeddings = []
3 labels = []
4 for data, label in reg_test_data:
5     _, embedding = reg_model(data, raw=True)
6     embeddings.append(embedding)
7     labels.append(label)
8 embeddings = torch.cat(embeddings, axis=0)
9 labels = torch.cat(labels, axis=0)
10
11 # Detach the tensor and convert to NumPy for PCA processing
12 if embeddings.requires_grad:
13     embeddings_np = embeddings.detach().cpu().numpy()
14 else:
15     embeddings_np = embeddings.cpu().numpy()
16
17 # Perform PCA to reduce to 2 dimensions for visualization
18 pca = PCA(n_components=2)
19 reduced_embeddings = pca.fit_transform(embeddings_np)
20
21 # Plotting the embeddings with labels as color codes
22 plt.figure(figsize=(10, 8))
23 scatter = plt.scatter(reduced_embeddings[:, 0], reduced_embeddings[:, 1], c=labels, cmap='viridis', alpha=0.6)
24
25 plt.colorbar(scatter) # Show color scale
26 plt.title('PCA of Regression Latent Vectors')
27 plt.xlabel('Principal Component 1')
28 plt.ylabel('Principal Component 2')
29 plt.grid(True)
30 plt.show()

```



✓ Interactive 3D projection of regression latent vectors